

eco482project (5) (2) (1) (2)

April 6, 2025

```
[1]: import pandas as pd
import numpy as np
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestClassifier
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import seaborn as sns
import statsmodels.api as sm
#Bias + how i rates j + own i rating + demographic of i + demographic of j
#Define y clearly in presentation
file = 'speed_dating_with_participant_ids.csv'
df = pd.read_csv(file)
df
```

/tmp/ipykernel_100/3825651112.py:20: DtypeWarning: Columns (4,11,12,40,41,42,74,75,76,77,78,79,80,81,82,83,84,85,86,87,88,89,90,110) have mixed types. Specify dtype option on import or set low_memory=False.

```
df = pd.read_csv(file)
```

```
[1]:
```

	id	has_null	wave	gender	age	age_o	d_age	d_d_age	\
0	8202	1	21	male	23	22	1	[0-1]	
1	8193	1	21	male	23	26	3	[2-3]	
2	8194	1	21	male	23	22	1	[0-1]	
3	8195	1	21	male	23	27	4	[4-6]	
4	8196	1	21	male	23	25	2	[2-3]	
...
8373	7494	1	21	female	?	30	30	[7-37]	
8374	7496	1	21	female	?	30	30	[7-37]	
8375	7497	1	21	female	?	27	27	[7-37]	

8376	7483	1	21	female	?	22	22	[7-37]
8377	7495	1	21	female	?	28	28	[7-37]

	race \
0	Asian/Pacific Islander/Asian-American
1	Asian/Pacific Islander/Asian-American
2	Asian/Pacific Islander/Asian-American
3	Asian/Pacific Islander/Asian-American
4	Asian/Pacific Islander/Asian-American
...	...
8373	European/Caucasian-American
8374	European/Caucasian-American
8375	European/Caucasian-American
8376	European/Caucasian-American
8377	European/Caucasian-American

	race_o	...	d_like	d_guess_prob_liked	\
0	Asian/Pacific Islander/Asian-American	...	[6-8]	[5-6]	
1	European/Caucasian-American	...	[6-8]	[5-6]	
2	European/Caucasian-American	...	[6-8]	[7-10]	
3	Asian/Pacific Islander/Asian-American	...	[6-8]	[7-10]	
4	Asian/Pacific Islander/Asian-American	...	[6-8]	[5-6]	
...	
8373	European/Caucasian-American	...	[0-5]	[0-4]	
8374	European/Caucasian-American	...	[6-8]	[7-10]	
8375	Black/African American	...	[0-5]	[5-6]	
8376	Asian/Pacific Islander/Asian-American	...	[6-8]	[7-10]	
8377	European/Caucasian-American	...	[0-5]	[5-6]	

	met	decision	decision_o	match	Unnamed: 124	Unnamed: 125	Unnamed: 126	\
0	0	1	1	1	NaN	NaN	NaN	
1	0	1	0	0	NaN	NaN	NaN	
2	0	1	1	1	NaN	NaN	NaN	
3	1	1	1	1	NaN	NaN	NaN	
4	0	0	1	0	NaN	NaN	NaN	
...	
8373	?	0	0	0	NaN	NaN	NaN	
8374	0	0	1	0	NaN	NaN	NaN	
8375	0	0	0	0	NaN	NaN	NaN	
8376	0	0	0	0	NaN	NaN	NaN	
8377	0	0	1	0	NaN	NaN	NaN	

	participant_id
0	0
1	0
2	0
3	0

```

4          0
...
8373      526
8374      526
8375      526
8376      526
8377      526

```

[8378 rows x 128 columns]

```

[2]: columns_to_check = [
    "age", "race", "field", "gender", "attractive_o", "sinsere_o",
    ↪ "intelligence_o",
    "funny_o", "ambitious_o", "attractive", "sincere", "intelligence", "funny",
    "ambition", "expected_num_matches", "decision_o", "guess_prob_liked",
    ↪ "d_age",
]

# Replace "?" with NaN
df = df.replace("?", pd.NA)

# Drop rows with NaN in any of the specified columns
df = df.dropna(subset=columns_to_check)

df = df[df["guess_prob_liked"] != 5]
df

```

```

[2]:      id  has_null  wave  gender  age  age_o  d_age  d_d_age  \
0      8202         1    21   male   23    22     1    [0-1]
3      8195         1    21   male   23    27     4    [4-6]
4      8196         1    21   male   23    25     2    [2-3]
5      8197         1    21   male   23    24     1    [0-1]
6      8198         1    21   male   23    26     3    [2-3]
...
8278  6827         1    18  female   55    33    22    [7-37]
8279  6826         1    18  female   55    23    32    [7-37]
8280  6825         1    18  female   55    33    22    [7-37]
8281  6823         1    18  female   55    33    22    [7-37]
8282  6824         1    18  female   55    27    28    [7-37]

      race  \
0  Asian/Pacific Islander/Asian-American
3  Asian/Pacific Islander/Asian-American
4  Asian/Pacific Islander/Asian-American
5  Asian/Pacific Islander/Asian-American
6  Asian/Pacific Islander/Asian-American
...

```

8278 Asian/Pacific Islander/Asian-American
 8279 Asian/Pacific Islander/Asian-American
 8280 Asian/Pacific Islander/Asian-American
 8281 Asian/Pacific Islander/Asian-American
 8282 Asian/Pacific Islander/Asian-American

	race_o	...	d_like	d_guess_prob_liked	\
0	Asian/Pacific Islander/Asian-American	...	[6-8]	[5-6]	
3	Asian/Pacific Islander/Asian-American	...	[6-8]	[7-10]	
4	Asian/Pacific Islander/Asian-American	...	[6-8]	[5-6]	
5	European/Caucasian-American	...	[6-8]	[0-4]	
6	Latino/Hispanic American	...	[0-5]	[0-4]	
...	
8278	Asian/Pacific Islander/Asian-American	...	[6-8]	[0-4]	
8279	Asian/Pacific Islander/Asian-American	...	[6-8]	[0-4]	
8280	Asian/Pacific Islander/Asian-American	...	[0-5]	[0-4]	
8281	European/Caucasian-American	...	[6-8]	[0-4]	
8282	Other	...	[0-5]	[0-4]	

	met	decision	decision_o	match	Unnamed: 124	Unnamed: 125	Unnamed: 126	\
0	0	1	1	1	NaN	NaN	NaN	
3	1	1	1	1	NaN	NaN	NaN	
4	0	0	1	0	NaN	NaN	NaN	
5	0	0	0	0	NaN	NaN	NaN	
6	0	0	0	0	NaN	NaN	NaN	
...	
8278	0	1	0	0	NaN	NaN	NaN	
8279	0	1	0	0	NaN	NaN	NaN	
8280	0	0	1	0	NaN	NaN	NaN	
8281	0	1	0	0	NaN	NaN	NaN	
8282	0	0	0	0	NaN	NaN	NaN	

	participant_id
0	0
3	0
4	0
5	0
6	0
...	...
8278	523
8279	523
8280	523
8281	523
8282	523

[6316 rows x 128 columns]

```
[3]: df["guess_prob_liked"] = pd.to_numeric(df["guess_prob_liked"], errors="coerce")
df["decision_o"] = pd.to_numeric(df["decision_o"], errors="coerce")
```

```
# Apply the logic using np.where
```

```
import numpy as np
```

```
df["prediction_like"] = np.where(
    ((df["guess_prob_liked"] > 5) & (df["decision_o"] == 0)) |
    ((df["guess_prob_liked"] < 5) & (df["decision_o"] == 1)),
    0, # Miscalibrated
    1 # Calibrated
)
```

```
[4]: df = df.loc[:, ~df.columns.str.contains('^Unnamed')]
```

```
[5]: characteristics = {
    "attractive": "attractive_o",
    "sincere": "sinsere_o",
    "intelligence": "intelligence_o",
    "funny": "funny_o",
    "ambition": "ambitious_o"
}

# Make sure all relevant columns are numeric
for col in list(characteristics.keys()) + list(characteristics.values()):
    df[col] = pd.to_numeric(df[col], errors='coerce')

# Dictionary to store the average bias per person per trait
participant_biases = {}

# Loop through each characteristic
for self_col, other_col in characteristics.items():
    # Calculate self - other rating (row-wise)
    df[f'diff_{self_col}'] = df[self_col] - df[other_col]

    # Compute average difference per participant
    avg_bias = df.groupby("participant_id")[f'diff_{self_col}'].mean()

    # Save the average bias
    df[f'bias_{self_col}'] = df["participant_id"].map(avg_bias)

# Drop intermediate columns if desired
df.drop(columns=[f'diff_{col}' for col in characteristics.keys()], inplace=True)

# Show the updated DataFrame
df.head()
```

```
/tmp/ipykernel_100/1577411713.py:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df[col] = pd.to_numeric(df[col], errors='coerce')
/tmp/ipykernel_100/1577411713.py:19: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df[f'diff_{self_col}'] = df[self_col] - df[other_col]
/tmp/ipykernel_100/1577411713.py:25: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df[f'bias_{self_col}'] = df["participant_id"].map(avg_bias)
```

```
[5]:
```

	id	has_null	wave	gender	age	age_o	d_age	d_d_age	\
0	8202	1	21	male	23	22	1	[0-1]	
3	8195	1	21	male	23	27	4	[4-6]	
4	8196	1	21	male	23	25	2	[2-3]	
5	8197	1	21	male	23	24	1	[0-1]	
6	8198	1	21	male	23	26	3	[2-3]	

	race	\
0	Asian/Pacific Islander/Asian-American	
3	Asian/Pacific Islander/Asian-American	
4	Asian/Pacific Islander/Asian-American	
5	Asian/Pacific Islander/Asian-American	
6	Asian/Pacific Islander/Asian-American	

	race_o	...	decision	decision_o	match	\
0	Asian/Pacific Islander/Asian-American	...	1	1	1	
3	Asian/Pacific Islander/Asian-American	...	1	1	1	
4	Asian/Pacific Islander/Asian-American	...	0	1	0	
5	European/Caucasian-American	...	0	0	0	
6	Latino/Hispanic American	...	0	0	0	

	participant_id	prediction_like	bias_attractive	bias_sincere	\
0	0	1	2.0	4.5	
3	0	1	2.0	4.5	
4	0	1	2.0	4.5	

5	0	1	2.0	4.5
6	0	1	2.0	4.5

	bias_intelligence	bias_funny	bias_ambition
0	4.0	3.75	2.25
3	4.0	3.75	2.25
4	4.0	3.75	2.25
5	4.0	3.75	2.25
6	4.0	3.75	2.25

[5 rows x 131 columns]

```
[6]: df["expected_num_matches"] = pd.to_numeric(df["expected_num_matches"],
        errors="coerce")
df["match"] = pd.to_numeric(df["match"], errors="coerce")

# Compute total actual matches per participant
total_matches = df.groupby("participant_id")["match"].sum()

# Compute expected matches per participant (mean across their rows)
expected_matches = df.groupby("participant_id")["expected_num_matches"].mean()

# Compute matching bias: actual - expected
matching_bias = total_matches - expected_matches

# Add matching bias to each row of the DataFrame
df["matching_bias"] = df["participant_id"].map(matching_bias)
df.head()
```

```
[6]:      id  has_null  wave gender age age_o  d_age d_d_age \
0  8202         1    21   male  23   22     1    [0-1]
3  8195         1    21   male  23   27     4    [4-6]
4  8196         1    21   male  23   25     2    [2-3]
5  8197         1    21   male  23   24     1    [0-1]
6  8198         1    21   male  23   26     3    [2-3]
```

	race
0	Asian/Pacific Islander/Asian-American
3	Asian/Pacific Islander/Asian-American
4	Asian/Pacific Islander/Asian-American
5	Asian/Pacific Islander/Asian-American
6	Asian/Pacific Islander/Asian-American

	race_o	...	decision_o	match
0	Asian/Pacific Islander/Asian-American	...	1	1
3	Asian/Pacific Islander/Asian-American	...	1	1
4	Asian/Pacific Islander/Asian-American	...	1	0

5	European/Caucasian-American ...	0	0
6	Latino/Hispanic American ...	0	0

	participant_id	prediction_like	bias_attractive	bias_sincere \
0	0	1	2.0	4.5
3	0	1	2.0	4.5
4	0	1	2.0	4.5
5	0	1	2.0	4.5
6	0	1	2.0	4.5

	bias_intelligence	bias_funny	bias_ambition	matching_bias
0	4.0	3.75	2.25	2.0
3	4.0	3.75	2.25	2.0
4	4.0	3.75	2.25	2.0
5	4.0	3.75	2.25	2.0
6	4.0	3.75	2.25	2.0

[5 rows x 132 columns]

```
[7]: df["field"] = df["field"].astype(str).str.strip().str.lower()
df["field_o"] = np.nan # initialize the new column

# Ensure age columns are numeric
df["age"] = pd.to_numeric(df["age"], errors="coerce")
df["age_o"] = pd.to_numeric(df["age_o"], errors="coerce")

# Define preference importance column mappings
preference_match_columns = {
    "pref_o_attractive": "attractive_important",
    "pref_o_sincere": "sincere_important",
    "pref_o_intelligence": "intelligence_important",
    "pref_o_funny": "funny_important",
    "pref_o_ambitious": "ambition_important",
    "pref_o_shared_interests": "shared_interests_important"
}

# Ensure all involved columns are numeric
for pref_col, imp_col in preference_match_columns.items():
    df[pref_col] = pd.to_numeric(df[pref_col], errors="coerce")
    df[imp_col] = pd.to_numeric(df[imp_col], errors="coerce")

df["same_field"] = 0

for wave in df["wave"].unique():
    wave_df = df[df["wave"] == wave]
    males = wave_df[wave_df["gender"] == "male"]
    females = wave_df[wave_df["gender"] == "female"]
```



```

for idx, male in males.iterrows():
    matches = females[
        (females["age"] == male["age_o"])
    ]
    for pref_col, imp_col in preference_match_columns.items():
        matches = matches[matches[imp_col] == male[pref_col]]

    if not matches.empty:
        # Choose the first matching partner
        matched_female = matches.iloc[0]
        if matched_female["field"] == male["field"]:
            df.loc[idx, "same_field"] = 1
            df.loc[idx, "field_o"] = matched_female["field"]

for idx, female in females.iterrows():
    matches = males[
        (males["age"] == female["age_o"])
    ]
    for pref_col, imp_col in preference_match_columns.items():
        matches = matches[matches[imp_col] == female[pref_col]]

    if not matches.empty:
        matched_male = matches.iloc[0]
        if matched_male["field"] == female["field"]:
            df.loc[idx, "same_field"] = 1
            df.loc[idx, "field_o"] = matched_male["field"]

print(df["same_field"].sum())
df.head()

```

/tmp/ipykernel_100/3114844432.py:42: FutureWarning: Setting an item of incompatible dtype is deprecated and will raise in a future error of pandas. Value 'mechanical engineering' has dtype incompatible with float64, please explicitly cast to a compatible dtype first.

```
df.loc[idx, "field_o"] = matched_female["field"]
```

152

```
[7]:
```

	id	has_null	wave	gender	age	age_o	d_age	d_d_age	\
0	8202	1	21	male	23	22.0	1	[0-1]	
3	8195	1	21	male	23	27.0	4	[4-6]	
4	8196	1	21	male	23	25.0	2	[2-3]	
5	8197	1	21	male	23	24.0	1	[0-1]	
6	8198	1	21	male	23	26.0	3	[2-3]	

```

race \
0 Asian/Pacific Islander/Asian-American

```

```

3 Asian/Pacific Islander/Asian-American
4 Asian/Pacific Islander/Asian-American
5 Asian/Pacific Islander/Asian-American
6 Asian/Pacific Islander/Asian-American

```

	race_o ...	participant_id	prediction_like \
0	Asian/Pacific Islander/Asian-American ...	0	1
3	Asian/Pacific Islander/Asian-American ...	0	1
4	Asian/Pacific Islander/Asian-American ...	0	1
5	European/Caucasian-American ...	0	1
6	Latino/Hispanic American ...	0	1

	bias_attractive	bias_sincere	bias_intelligence	bias_funny	bias_ambition \
0	2.0	4.5	4.0	3.75	2.25
3	2.0	4.5	4.0	3.75	2.25
4	2.0	4.5	4.0	3.75	2.25
5	2.0	4.5	4.0	3.75	2.25
6	2.0	4.5	4.0	3.75	2.25

	matching_bias	field_o	same_field
0	2.0	mechanical engineering	0
3	2.0	NaN	0
4	2.0	medicine	0
5	2.0	public health	0
6	2.0	social work	0

[5 rows x 134 columns]

```
[8]: df["gender_o"] = df["gender"].apply(lambda g: "female" if g == "male" else
    ↪ "male")
```

```
[9]: # List of predictors: d_age, samerace, interests_correlate, same_field,
    ↪ bias_attractive, bias_sincere, bias_intelligence, bias_funny, bias_ambition,
    ↪ matching_bias, gender, race, age, field
```

```
[10]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import ttest_ind
import numpy as np

# -----
# Step 0: Prepare Encodings for Gender, Race, Field
# -----
df['gender_num'] = df['gender'].str.lower().map({'male': 1, 'female': 0})
df['gender_num_o'] = df['gender_o'].str.lower().map({'male': 1, 'female': 0})
df['race_num'] = df['race'].astype('category').cat.codes
```

```

df['race_o_num'] = df['race_o'].astype('category').cat.codes

all_fields = pd.concat([df["field"], df["field_o"]], ignore_index=True).
    ↪astype('category')
consistent_categories = all_fields.cat.categories
df["field"] = pd.Categorical(df["field"], categories=consistent_categories)
df["field_o"] = pd.Categorical(df["field_o"], categories=consistent_categories)
df["field_num"] = df["field"].cat.codes
df["field_num_o"] = df["field_o"].cat.codes

# -----
# Step 1: Convert other key columns to numeric
# -----
numeric_cols = [
    "decision", "decision_o", "guess_prob_liked", "like",
    "expected_num_matches", "matching_bias",
    "bias_attractive", "bias_sincere", "bias_intelligence",
    "bias_funny", "bias_ambition",
    "attractive", "sincere", "intelligence", "funny", "ambition"
]
df[numeric_cols] = df[numeric_cols].apply(pd.to_numeric, errors="coerce")

# Asterisk function
def p_to_stars(p):
    if p < 0.001:
        return '***'
    elif p < 0.01:
        return '**'
    elif p < 0.05:
        return '*'
    else:
        return ''

# Function to compute mean and SE
def mean_and_se(df, group_col, target_cols):
    means = df.groupby(group_col)[target_cols].mean()
    stds = df.groupby(group_col)[target_cols].std()
    counts = df.groupby(group_col)[target_cols].count()
    ses = stds / np.sqrt(counts)
    return means.T, ses.T

# -----
# Step 2: Correlation Analysis
# -----
print("Correlation between guess_prob_liked and decision_o:",
    ↪df["guess_prob_liked"].corr(df["decision_o"]))
print("Correlation between like and decision:", df["like"].corr(df["decision"]))

```

```

# -----
# Step 3: Trait Ratings by Gender with Asterisks + SE
# -----

traits = ["attractive", "sincere", "intelligence", "funny", "ambition"]
trait_names = ["Attractiveness", "Sincerity", "Intelligence", "Humor", "Ambition"]

trait_means, trait_ses = mean_and_se(df, "gender_num", traits)
trait_means.index = trait_names
trait_ses.index = trait_names

# T-tests
pvals = []
for trait in traits:
    f = df[df["gender_num"] == 0][trait].dropna()
    m = df[df["gender_num"] == 1][trait].dropna()
    _, p = ttest_ind(f, m, equal_var=False)
    pvals.append(p_to_stars(p))

fig, ax = plt.subplots(figsize=(10, 6))
x = np.arange(len(trait_names))
bar_width = 0.35

# Bar plots with error bars
ax.bar(x - bar_width/2, trait_means[0].values, yerr=trait_ses[0].values, width=bar_width, label="Female", capsize=5)
ax.bar(x + bar_width/2, trait_means[1].values, yerr=trait_ses[1].values, width=bar_width, label="Male", capsize=5)

for i, stars in enumerate(pvals):
    if stars:
        ymax = max(trait_means.iloc[i][0] + trait_ses.iloc[i][0], trait_means.
        iloc[i][1] + trait_ses.iloc[i][1])
        ax.text(i, ymax + 0.1, stars, ha='center', color='red', fontsize=14)

ax.set_xticks(x)
ax.set_xticklabels(trait_names, rotation=45)
ax.set_title("Average Trait Ratings by Gender")
ax.set_ylabel("Average Rating")
ax.set_xlabel("Trait")
ax.legend(title="Gender")
plt.tight_layout()
plt.grid(False)
plt.show()

# -----

```

```

# Step 4: Expected Matches by Gender
# -----
group0 = df[df["gender_num"] == 0]["expected_num_matches"].dropna()
group1 = df[df["gender_num"] == 1]["expected_num_matches"].dropna()
_, p = ttest_ind(group0, group1, equal_var=False)
asterisks = p_to_stars(p)

means = [group0.mean(), group1.mean()]
ses = [group0.std() / np.sqrt(len(group0)), group1.std() / np.sqrt(len(group1))]

plt.figure(figsize=(6, 4))
ax = plt.gca()
ax.bar(["Female", "Male"], means, yerr=ses, capsize=5)
plt.title("Expected Matches by Gender")
plt.ylabel("Average Expected Matches")
plt.xlabel("Gender")
plt.tight_layout()
plt.grid(False)

if asterisks:
    ymax = max(means[i] + ses[i] for i in range(2))
    ax.text(0.5, ymax + 0.1, asterisks, ha='center', fontsize=14, color='red')

plt.show()

# -----
# Step 5: Self-Other Bias by Gender with SE
# -----
bias_cols = ["bias_attractive", "bias_sincere", "bias_intelligence",
             ↪ "bias_funny", "bias_ambition"]
bias_names = ["Attractiveness Bias", "Sincerity Bias", "Intelligence Bias",
             ↪ "Humor Bias", "Ambition Bias"]

bias_means, bias_ses = mean_and_se(df, "gender_num", bias_cols)
bias_means.index = bias_names
bias_ses.index = bias_names

pvals_bias = []
for col in bias_cols:
    f = df[df["gender_num"] == 0][col].dropna()
    m = df[df["gender_num"] == 1][col].dropna()
    _, p = ttest_ind(f, m, equal_var=False)
    pvals_bias.append(p_to_stars(p))

fig, ax = plt.subplots(figsize=(10, 6))
x = np.arange(len(bias_names))

```

```

ax.bar(x - bar_width/2, bias_means[0].values, yerr=bias_ses[0].values,
      width=bar_width, label="Female", capsize=5)
ax.bar(x + bar_width/2, bias_means[1].values, yerr=bias_ses[1].values,
      width=bar_width, label="Male", capsize=5)

for i, stars in enumerate(pvals_bias):
    if stars:
        ymax = max(bias_means.iloc[i][0] + bias_ses.iloc[i][0], bias_means.
      iloc[i][1] + bias_ses.iloc[i][1])
        ax.text(i, ymax + 0.1, stars, ha='center', color='red', fontsize=14)

ax.set_xticks(x)
ax.set_xticklabels(bias_names, rotation=45)
ax.set_title("Average Bias by Characteristic and Gender")
ax.set_ylabel("Average Bias (Self - Other)")
ax.set_xlabel("Characteristic")
ax.legend(title="Gender")
plt.tight_layout()
plt.grid(False)
plt.show()

# -----
# Step 6: Matching Bias by Gender with SE
# -----
g0 = df[df["gender_num"] == 0]["matching_bias"].dropna()
g1 = df[df["gender_num"] == 1]["matching_bias"].dropna()
_, p = ttest_ind(g0, g1, equal_var=False)
asterisks = p_to_stars(p)

means = [g0.mean(), g1.mean()]
ses = [g0.std() / np.sqrt(len(g0)), g1.std() / np.sqrt(len(g1))]


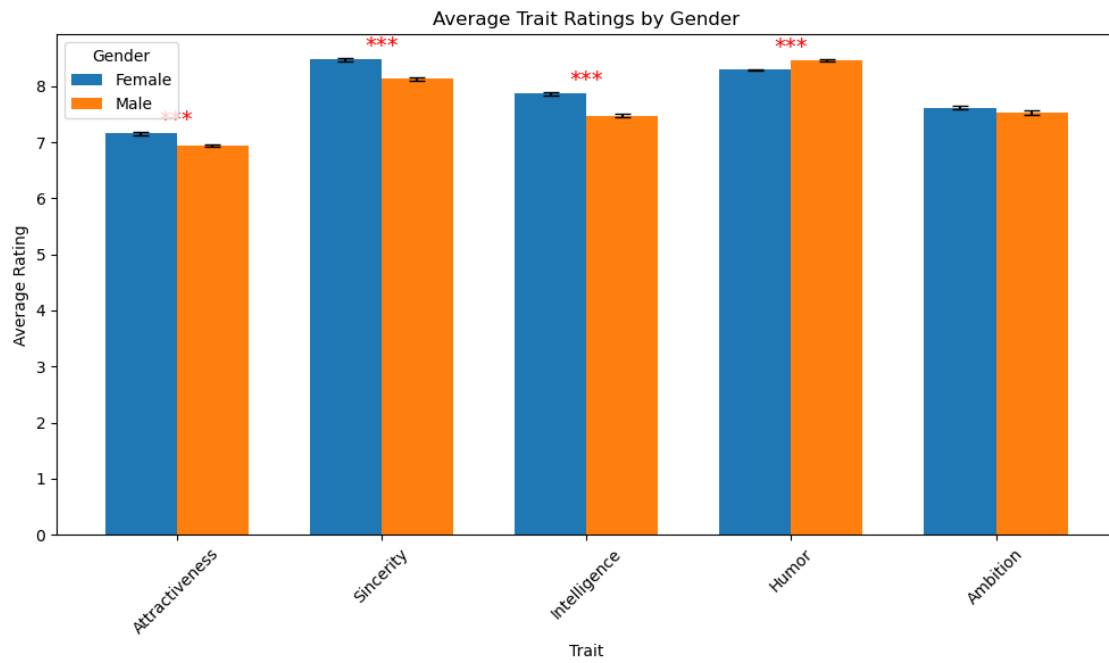
plt.figure(figsize=(6, 4))
ax = plt.gca()
ax.bar(["Female", "Male"], means, yerr=ses, capsize=5)
plt.title("Matching Bias by Gender")
plt.xlabel("Gender")
plt.ylabel("Matching Bias (Actual - Expected Matches)")
plt.tight_layout()
plt.grid(False)

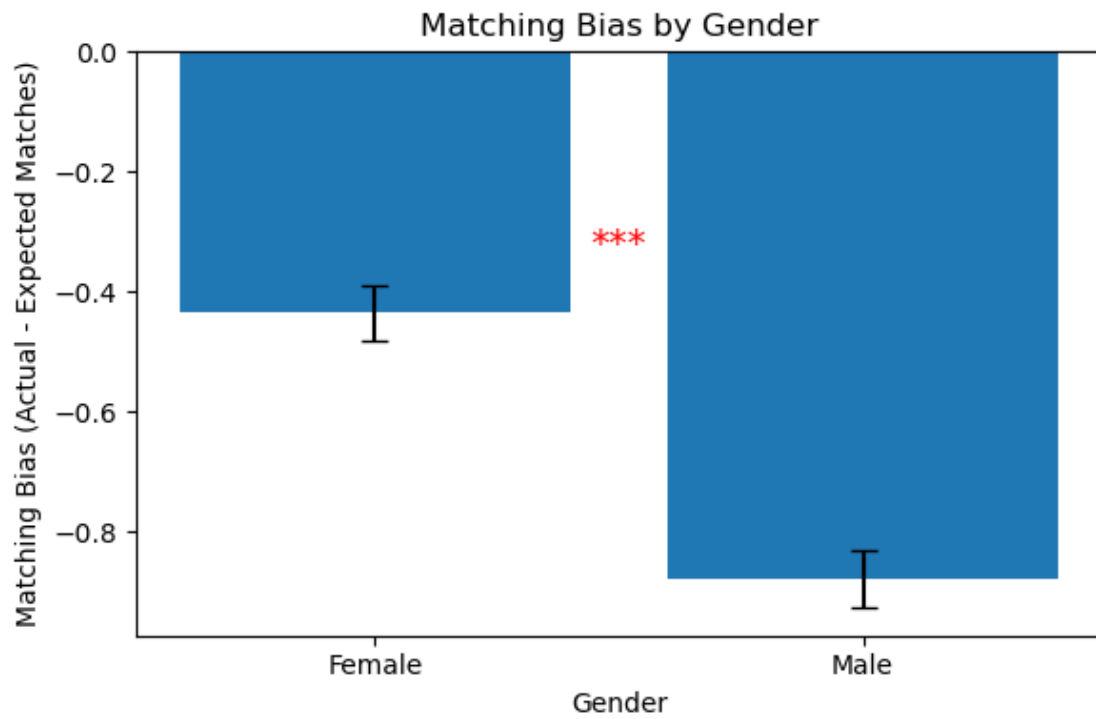
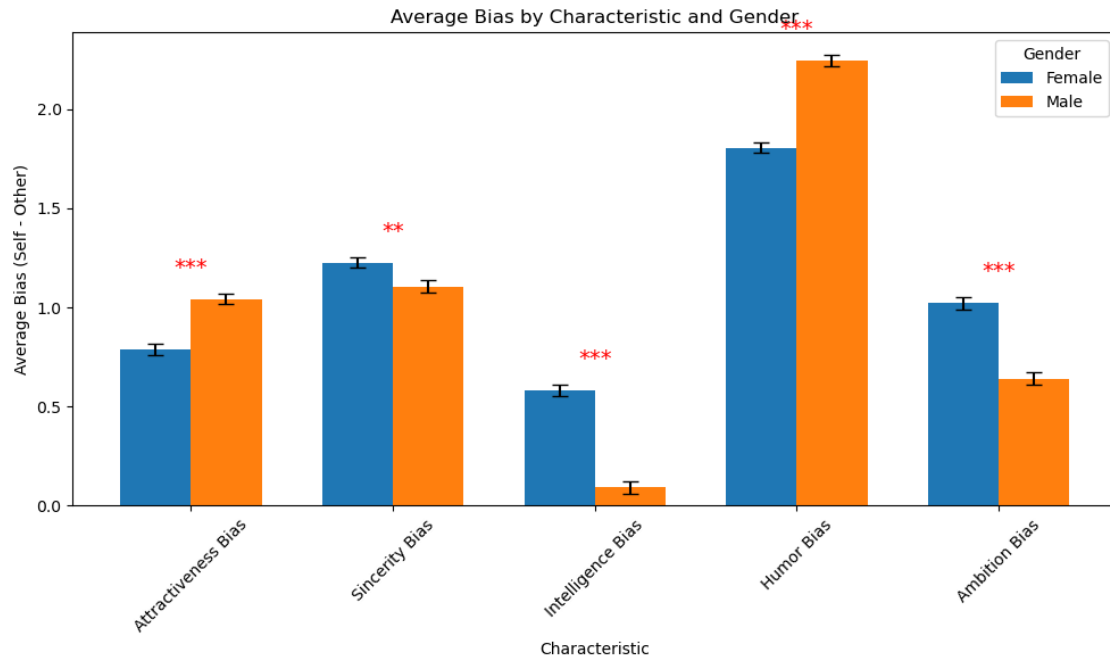
if asterisks:
    ymax = max(means[i] + ses[i] for i in range(2))
    ax.text(0.5, ymax + 0.05, asterisks, ha='center', color='red', fontsize=14)

plt.show()

```

Correlation between guess_prob_liked and decision_o: 0.1409880632116935
Correlation between like and decision: 0.5150175200729944






```

[11]: import matplotlib.pyplot as plt
import scipy.stats as stats
import numpy as np
import pandas as pd

# Clean gender column
df['gender'] = df['gender'].str.lower().str.strip()

# Compute group means and standard errors
group_stats = df.groupby('gender')['prediction_like'].agg(['mean', 'count',
↳ 'std'])
group_stats['sem'] = group_stats['std'] / np.sqrt(group_stats['count'])

# Perform t-test between gender groups (assuming two groups)
groups = df['gender'].unique()
if len(groups) == 2:
    group1 = df[df['gender'] == groups[0]]['decision']
    group2 = df[df['gender'] == groups[1]]['decision']
    t_stat, p_value = stats.ttest_ind(group1, group2, equal_var=False)

    # Determine significance level
    if p_value < 0.001:
        sig_label = '***'
    elif p_value < 0.01:
        sig_label = '**'
    elif p_value < 0.05:
        sig_label = '*'
    else:
        sig_label = 'n.s.' # Not significant
else:
    sig_label = ''
    print("Warning: More than two gender groups, skipping t-test.")

# Plot with error bars
plt.figure(figsize=(6, 4))
bars = plt.bar(group_stats.index, group_stats['mean'],
               yerr=group_stats['sem'], capsize=5,
               color=['skyblue', 'salmon'])

plt.title("Average Prediction_Like by Gender")
plt.ylabel("Average Prediction_Like")
plt.xlabel("Gender")
plt.ylim(0, 1)
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)

# Annotate with significance stars


```

```

if sig_label and sig_label != 'n.s.':
    max_height = group_stats['mean'].max() + group_stats['sem'].max() + 0.02
    plt.text(0.5, max_height, sig_label, ha='center', fontsize=14)

plt.tight_layout()
plt.show()

```



```

[12]: import pandas as pd
import numpy as np
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import classification_report
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

# Encode categorical variables
df['gender_num'] = df['gender'].str.lower().map({'male': 1, 'female': 0})
df['gender_num_o'] = df['gender_o'].str.lower().map({'male': 1, 'female': 0})
df['race_num'] = df['race'].astype('category').cat.codes
df['race_o_num'] = df['race_o'].astype('category').cat.codes
all_fields = pd.concat([df["field"], df["field_o"]], ignore_index=True).
    .astype('category')
consistent_categories = all_fields.cat.categories

```

```

# Apply the same categories to both columns
df["field"] = pd.Categorical(df["field"], categories=consistent_categories)
df["field_o"] = pd.Categorical(df["field_o"], categories=consistent_categories)

# Now generate codes
df["field_num"] = df["field"].cat.codes
df["field_num_o"] = df["field_o"].cat.codes

# Define feature columns (include newly added partner-related variables)
features = [
    'd_age', 'samerace', 'interests_correlate', 'same_field',
    'bias_attractive', 'bias_sincere', 'bias_intelligence',
    'bias_funny', 'bias_ambition',
    'gender_num', 'race_num', 'age', 'field_num',
    'race_o_num', 'age_o',
    'attractive_partner', 'sincere_partner', 'funny_partner',
    'intelligence_partner', 'ambition_partner', 'gender_num_o',
    'field_num_o', 'attractive', 'sincere', 'intelligence', 'funny',
    'ambition'
]

# Prepare feature matrix and target
X = df[features]
y = pd.to_numeric(df['prediction_like'], errors='coerce')

# Drop rows with missing values
model_df = pd.concat([X, y], axis=1).dropna()
X_clean = model_df[features]
y_clean = model_df['prediction_like']

# Use cross-validation to find best k
k_range = range(1, int(np.sqrt(len(X_clean))) + 1)
cv_scores = []

for k in k_range:
    knn_pipeline = Pipeline([
        ('scaler', StandardScaler()),
        ('knn', KNeighborsClassifier(n_neighbors=k))
    ])
    scores = cross_val_score(knn_pipeline, X_clean, y_clean, cv=5,
        scoring='accuracy')
    cv_scores.append(scores.mean())

best_k = k_range[np.argmax(cv_scores)]
print(f"Best k: {best_k} with cross-validated accuracy: {max(cv_scores):.4f}")

```

```

# Final train/test split and model fit
X_train, X_test, y_train, y_test = train_test_split(X_clean, y_clean,
    ↪test_size=0.2, random_state=42)

final_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('knn', KNeighborsClassifier(n_neighbors=best_k))
])
final_pipeline.fit(X_train, y_train)
y_pred = final_pipeline.predict(X_test)

print("\n=== Final KNN Classification Report ===")
print(classification_report(y_test, y_pred))

```

Best k: 71 with cross-validated accuracy: 0.6483

```

=== Final KNN Classification Report ===

```

	precision	recall	f1-score	support
0	0.79	0.04	0.07	399
1	0.66	0.99	0.79	751
accuracy			0.66	1150
macro avg	0.72	0.52	0.43	1150
weighted avg	0.71	0.66	0.54	1150

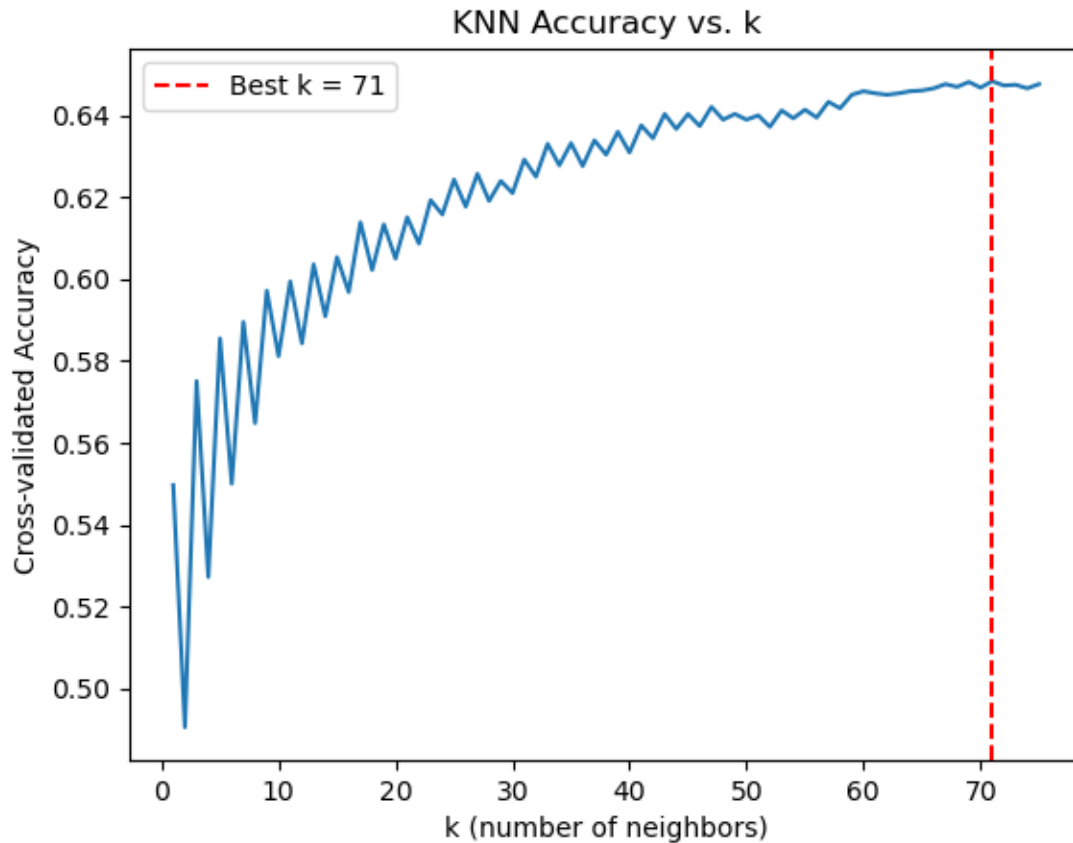
```

[13]: import matplotlib.pyplot as plt

plt.plot(k_range, cv_scores)
plt.xlabel("k (number of neighbors)")
plt.ylabel("Cross-validated Accuracy")
plt.title("KNN Accuracy vs. k")
plt.axvline(x=best_k, color='red', linestyle='--', label=f"Best k = {best_k}")
plt.legend()
plt.show()

#The best value of k is 75, but the value of a higher k begins to diminish
    ↪around k=40. To avoid underfitting and oversmoothing, we pick k=40

```



```
[14]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import classification_report, confusion_matrix, \
    ConfusionMatrixDisplay, roc_auc_score, roc_curve
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

# === Encode Categorical Variables ===
df['gender_num'] = df['gender'].str.lower().map({'male': 1, 'female': 0})
df['gender_num_o'] = df['gender_o'].str.lower().map({'male': 1, 'female': 0})
df['race_num'] = df['race'].astype('category').cat.codes
df['race_o_num'] = df['race_o'].astype('category').cat.codes

all_fields = pd.concat([df["field"], df["field_o"]], ignore_index=True).
    .astype('category')
consistent_categories = all_fields.cat.categories
df["field"] = pd.Categorical(df["field"], categories=consistent_categories)
```

```

df["field_o"] = pd.Categorical(df["field_o"], categories=consistent_categories)
df["field_num"] = df["field"].cat.codes
df["field_num_o"] = df["field_o"].cat.codes

# === Define Features and Target ===
features = [
    'd_age', 'samerace', 'interests_correlate', 'same_field',
    'bias_attractive', 'bias_sincere', 'bias_intelligence',
    'bias_funny', 'bias_ambition',
    'gender_num', 'race_num', 'age', 'field_num',
    'race_o_num', 'age_o',
    'attractive_partner', 'sincere_partner', 'funny_partner',
    'intelligence_partner', 'ambition_partner', 'gender_num_o',
    'field_num_o', 'attractive', 'sincere', 'intelligence', 'funny',
    'ambition'
]

X = df[features]
y = pd.to_numeric(df['prediction_like'], errors='coerce')

# === Drop missing values ===
model_df = pd.concat([X, y], axis=1).dropna()
X_clean = model_df[features]
y_clean = model_df['prediction_like']

# === Train-Test Split ===
X_train, X_test, y_train, y_test = train_test_split(X_clean, y_clean,
                                                    test_size=0.2, random_state=42)

# === Define and Fit KNN Pipeline (k=40) ===
knn_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('knn', KNeighborsClassifier(n_neighbors=40))
])
knn_pipeline.fit(X_train, y_train)
y_pred = knn_pipeline.predict(X_test)
y_prob = knn_pipeline.predict_proba(X_test)[:, 1]

# === Evaluation ===
print("\n=== Final KNN Classification Report (k=40) ===")
print(classification_report(y_test, y_pred))

# === Confusion Matrix ===
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap="Blues")
plt.title("Confusion Matrix - KNN (k = 40)")

```

```

plt.show()

# === ROC Curve and AUC ===
fpr, tpr, _ = roc_curve(y_test, y_prob)
auc = roc_auc_score(y_test, y_prob)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {auc:.2f})', linewidth=2)
plt.plot([0, 1], [0, 1], 'k--', label='Random Guess')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - KNN (k = 40)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

print(f"AUC Score: {auc:.4f}")

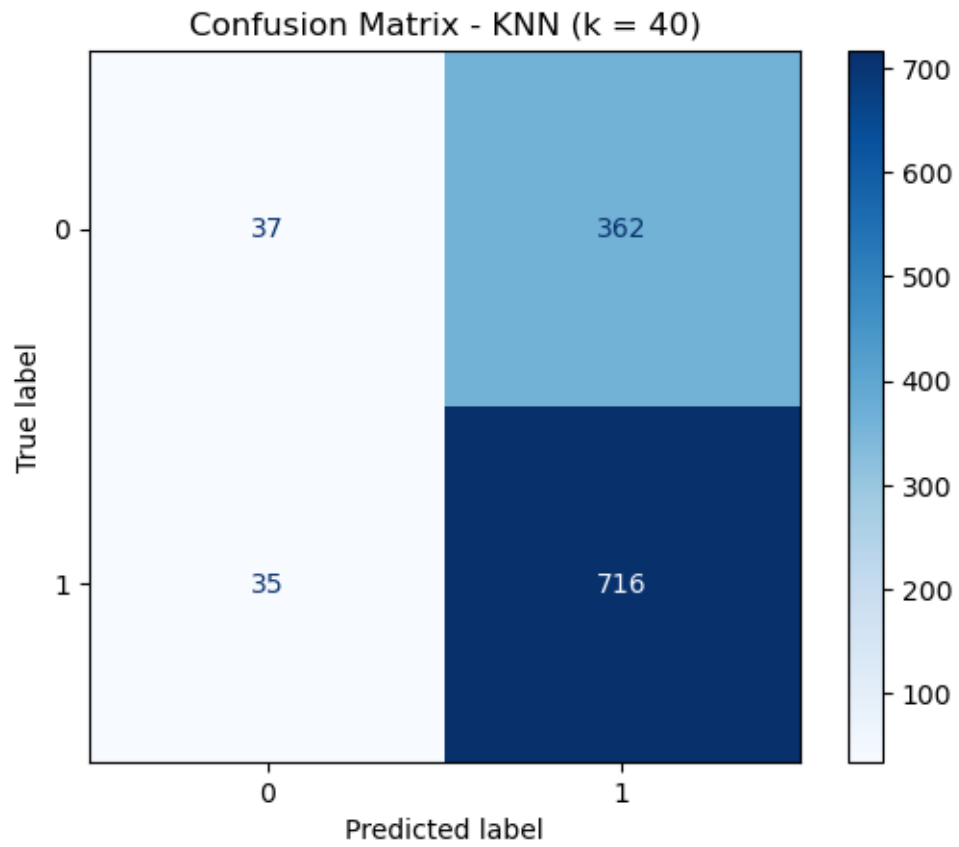
```


```

=== Final KNN Classification Report (k=40) ===

```

	precision	recall	f1-score	support
0	0.51	0.09	0.16	399
1	0.66	0.95	0.78	751
accuracy			0.65	1150
macro avg	0.59	0.52	0.47	1150
weighted avg	0.61	0.65	0.57	1150





AUC Score: 0.5803

```
[15]: #cross validation for KNN
cv_scores = cross_val_score(knn_pipeline, X_clean, y_clean, cv=5,
                             scoring='accuracy')

print("Cross-validated scores (accuracy):", cv_scores)
print("Mean accuracy:", cv_scores.mean())
#KNN model's performance is pretty consistent, with ~65%-68% accuracy whether
#you do one train-test split or 5-fold CV
```

Cross-validated scores (accuracy): [0.64608696 0.63565217 0.63913043 0.63391304 0.59965187]

Mean accuracy: 0.6308868959775986

```
[16]: from sklearn.metrics import confusion_matrix, recall_score, mean_squared_error
from sklearn.model_selection import cross_val_score

# === False Positive Rate (FPR) ===
cm = confusion_matrix(y_test, y_pred)
```

```

tn, fp, fn, tp = cm.ravel()
fpr = fp / (fp + tn)
print(f"False Positive Rate (FPR): {fpr:.4f}")

# === Cross-Validated MSE ===
mse_scores = cross_val_score(final_pipeline, X_clean, y_clean, cv=5,
    ↪scoring='neg_mean_squared_error')
mean_mse = -mse_scores.mean()
print(f"Cross-Validated Mean Squared Error (MSE): {mean_mse:.4f}")

# === Recall Score ===
recall = recall_score(y_test, y_pred)
print(f"Recall Score: {recall:.4f}")

```

False Positive Rate (FPR): 0.9073
 Cross-Validated Mean Squared Error (MSE): 0.3517
 Recall Score: 0.9534

```

[17]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import (
    classification_report, confusion_matrix, ConfusionMatrixDisplay,
    roc_auc_score, roc_curve
)
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

# === STEP 1: Encode categorical variables ===
df['gender_num'] = df['gender'].str.lower().map({'male': 1, 'female': 0})
df['gender_num_o'] = df['gender_o'].str.lower().map({'male': 1, 'female': 0})
df['race_num'] = df['race'].astype('category').cat.codes
df['race_o_num'] = df['race_o'].astype('category').cat.codes
all_fields = pd.concat([df["field"], df["field_o"]], ignore_index=True).
    ↪astype('category')
df["field"] = pd.Categorical(df["field"], categories=all_fields.cat.categories)
df["field_o"] = pd.Categorical(df["field_o"], categories=all_fields.cat.
    ↪categories)
df["field_num"] = df["field"].cat.codes
df["field_num_o"] = df["field_o"].cat.codes

# === STEP 2: Define predictors and target ===
features = [
    'd_age', 'samerace', 'interests_correlate', 'same_field',
    'bias_attractive', 'bias_sincere', 'bias_intelligence',

```

```

    'bias_funny', 'bias_ambition',
    'gender_num', 'race_num', 'age', 'field_num',
    'race_o_num', 'age_o',
    'attractive_partner', 'sincere_partner', 'funny_partner',
    'intelligence_partner', 'ambition_partner', 'gender_num_o',
    'field_num_o', 'attractive', 'sincere', 'intelligence', 'funny',
    'ambition'
]
X = df[features]
y = pd.to_numeric(df['prediction_like'], errors='coerce')

# === STEP 3: Drop missing values ===
model_df = pd.concat([X, y], axis=1).dropna()
X_clean = model_df[features]
y_clean = model_df['prediction_like']

# === STEP 4: Train-test split ===
X_train, X_test, y_train, y_test = train_test_split(X_clean, y_clean,
    ↪test_size=0.2, random_state=42)

# === STEP 5: Lasso Logistic Regression Pipeline ===
lasso_pipe = Pipeline([
    ('scaler', StandardScaler()),
    ('logreg', LogisticRegression(penalty='l1', solver='liblinear',
    ↪max_iter=1000, class_weight='balanced'))
])

# === STEP 6: Hyperparameter tuning ===
param_grid = {'logreg__C': np.logspace(-4, 4, 20)}
grid = GridSearchCV(lasso_pipe, param_grid, cv=5, scoring='accuracy')
grid.fit(X_train, y_train)

# === STEP 7: Predict and Threshold Tuning ===
y_prob = grid.predict_proba(X_test)[:, 1]

# Adjust threshold (try different cutoffs like 0.5, 0.6, 0.7)
threshold = 0.6
y_pred = (y_prob > threshold).astype(int)

# === STEP 8: Evaluation ===
print(f"Best C: {grid.best_params_['logreg__C']}")
print(f"\n=== Lasso Logistic Regression Classification Report (Threshold =
    ↪{threshold}) ===")
print(classification_report(y_test, y_pred))

# === Confusion Matrix ===
cm = confusion_matrix(y_test, y_pred)

```

```

disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=grid.classes_)
disp.plot(cmap='Blues')
plt.title(f"Confusion Matrix - Lasso (Threshold = {threshold})")
plt.show()

# === ROC Curve and AUC ===
fpr, tpr, _ = roc_curve(y_test, y_prob)
auc = roc_auc_score(y_test, y_prob)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {auc:.2f})', linewidth=2)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Lasso Logistic Regression')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

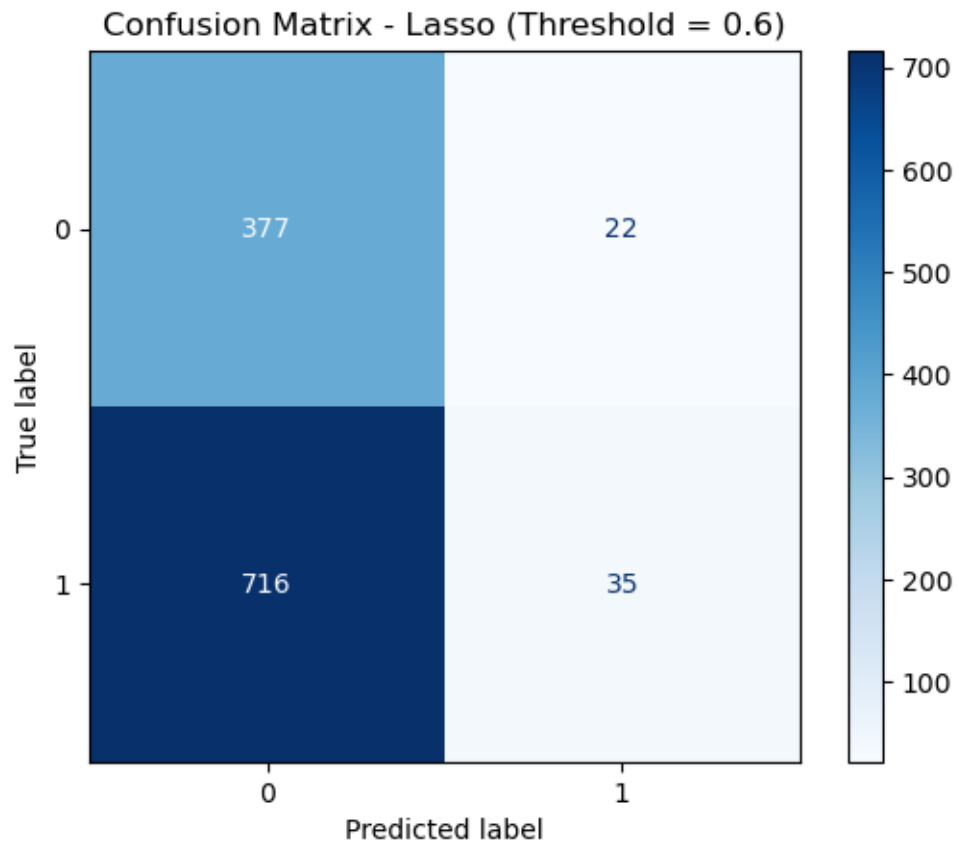
print(f"AUC Score: {auc:.4f}")


```

Best C: 11.288378916846883

=== Lasso Logistic Regression Classification Report (Threshold = 0.6) ===

	precision	recall	f1-score	support
0	0.34	0.94	0.51	399
1	0.61	0.05	0.09	751
accuracy			0.36	1150
macro avg	0.48	0.50	0.30	1150
weighted avg	0.52	0.36	0.23	1150





AUC Score: 0.5817

```
[18]: from sklearn.metrics import recall_score, confusion_matrix, make_scorer
      from sklearn.model_selection import cross_val_score
      from sklearn.metrics import mean_squared_error

      # === False Positive Rate ===
      tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
      false_positive_rate = fp / (fp + tn)
      print(f"False Positive Rate: {false_positive_rate:.4f}")

      # === Cross-validated Negative MSE ===
      neg_mse_scores = cross_val_score(
          grid.best_estimator_, X_clean, y_clean,
          scoring='neg_mean_squared_error', cv=5
      )
      mean_neg_mse = neg_mse_scores.mean()
      print(f"Cross-Validated Score (Negative MSE): {mean_neg_mse:.4f}")

      # === Recall Rate ===
```

```
recall = recall_score(y_test, y_pred)
print(f"Recall Rate: {recall:.4f}")
```

False Positive Rate: 0.0551
 Cross-Validated Score (Negative MSE): -0.4905
 Recall Rate: 0.0466

```
[19]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import classification_report, confusion_matrix, \
    ConfusionMatrixDisplay, roc_auc_score, roc_curve
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

# === STEP 1: Encode categorical variables ===
df['gender_num'] = df['gender'].str.lower().map({'male': 1, 'female': 0})
df['gender_num_o'] = df['gender_o'].str.lower().map({'male': 1, 'female': 0})
df['race_num'] = df['race'].astype('category').cat.codes
df['race_o_num'] = df['race_o'].astype('category').cat.codes
all_fields = pd.concat([df["field"], df["field_o"]], ignore_index=True).
    astype('category')
consistent_categories = all_fields.cat.categories
df["field"] = pd.Categorical(df["field"], categories=consistent_categories)
df["field_o"] = pd.Categorical(df["field_o"], categories=consistent_categories)
df["field_num"] = df["field"].cat.codes
df["field_num_o"] = df["field_o"].cat.codes

# === STEP 2: Define predictors and target ===
features = [
    'd_age', 'samerace', 'interests_correlate', 'same_field',
    'bias_attractive', 'bias_sincere', 'bias_intelligence',
    'bias_funny', 'bias_ambition',
    'gender_num', 'race_num', 'age', 'field_num',
    'race_o_num', 'age_o',
    'attractive_partner', 'sincere_partner', 'funny_partner',
    'intelligence_partner', 'ambition_partner', 'gender_num_o',
    'field_num_o', 'attractive', 'sincere', 'intelligence', 'funny',
    'ambition'
]
X = df[features]
y = pd.to_numeric(df['prediction_like'], errors='coerce')

# === STEP 3: Drop missing values ===
model_df = pd.concat([X, y], axis=1).dropna()
```

```

X_clean = model_df[features]
y_clean = model_df['prediction_like']

# === STEP 4: Train-test split ===
X_train, X_test, y_train, y_test = train_test_split(X_clean, y_clean,
    ↪test_size=0.2, random_state=42)

# === STEP 5: Ridge Logistic Regression Pipeline with Class Weights ===
ridge_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('logreg', LogisticRegression(
        penalty='l2',
        solver='liblinear',
        max_iter=1000,
        class_weight='balanced'
    ))
])

# === STEP 6: Hyperparameter tuning ===
param_grid = {'logreg__C': np.logspace(-4, 4, 20)}
grid = GridSearchCV(ridge_pipeline, param_grid, cv=5, scoring='accuracy')
grid.fit(X_train, y_train)

# === STEP 7: Predictions ===
y_pred = grid.predict(X_test)
y_prob = grid.predict_proba(X_test)[:, 1]

# === STEP 8: Evaluation Metrics ===
print(f"Best C (Regularization Strength): {grid.best_params_['logreg__C']}")
print("\n=== Ridge Logistic Regression (Balanced) Classification Report ===")
print(classification_report(y_test, y_pred))

# === STEP 9: Confusion Matrix ===
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=grid.classes_)
disp.plot(cmap='Blues')
plt.title("Confusion Matrix - Ridge Logistic Regression (Balanced)")
plt.show()

# === STEP 10: AUC and ROC Curve ===
fpr, tpr, _ = roc_curve(y_test, y_prob)
auc = roc_auc_score(y_test, y_prob)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {auc:.2f})', linewidth=2)
plt.plot([0, 1], [0, 1], 'k--', label='Random Guess')
plt.xlabel('False Positive Rate')

```



```

plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Ridge Logistic Regression (Balanced)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

print(f"\nAUC Score: {auc:.4f}")

# === STEP 11: Coefficients Summary ===
coefs = grid.best_estimator_.named_steps['logreg'].coef_[0]
coef_df = pd.DataFrame({
    'Feature': features,
    'Coefficient': coefs,
    'Odds Ratio': np.exp(coefs)
})
coef_df = coef_df.reindex(coef_df['Coefficient'].abs().
    ↪sort_values(ascending=False).index)

print("\nTop predictors (by absolute coefficient value):")
print(coef_df)


```


Best C (Regularization Strength): 0.615848211066026

=== Ridge Logistic Regression (Balanced) Classification Report ===

	precision	recall	f1-score	support
0	0.40	0.57	0.47	399
1	0.70	0.54	0.61	751
accuracy			0.55	1150
macro avg	0.55	0.55	0.54	1150
weighted avg	0.59	0.55	0.56	1150

Confusion Matrix - Ridge Logistic Regression (Balanced)





AUC Score: 0.5818

Top predictors (by absolute coefficient value):

	Feature	Coefficient	Odds Ratio
6	bias_intelligence	-0.240615	0.786144
5	bias_sincere	0.189271	1.208368
23	sincere	-0.164546	0.848278
7	bias_funny	-0.161962	0.850474
24	intelligence	0.151954	1.164107
25	funny	0.147839	1.159326
8	bias_ambition	0.107552	1.113548
26	ambition	-0.102188	0.902860
4	bias_attractive	-0.076819	0.926058
2	interests_correlate	0.065205	1.067377
18	intelligence_partner	-0.054384	0.947068
12	field_num	-0.053281	0.948113
16	sincere_partner	-0.049120	0.952067
10	race_num	-0.036928	0.963745
14	age_o	-0.033391	0.967160

15	attractive_partner	-0.032646	0.967881
20	gender_num_o	0.029064	1.029491
9	gender_num	-0.029064	0.971354
0	d_age	-0.029041	0.971377
1	samerace	-0.028815	0.971597
21	field_num_o	-0.026908	0.973451
19	ambition_partner	-0.024227	0.976064
13	race_o_num	-0.013148	0.986938
3	same_field	-0.010577	0.989479
22	attractive	-0.003882	0.996126
11	age	0.002006	1.002008
17	funny_partner	0.000114	1.000114

```
[20]: from sklearn.metrics import recall_score, confusion_matrix, make_scorer
from sklearn.model_selection import cross_val_score

# === False Positive Rate ===
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
false_positive_rate = fp / (fp + tn)
print(f"\nFalse Positive Rate: {false_positive_rate:.4f}")

# === Cross-Validated Negative MSE ===
neg_mse_scores = cross_val_score(
    grid.best_estimator_, X_clean, y_clean,
    scoring='neg_mean_squared_error', cv=5
)
mean_neg_mse = neg_mse_scores.mean()
print(f"Cross-Validated Score (Negative MSE): {mean_neg_mse:.4f}")

# === Recall Rate ===
recall = recall_score(y_test, y_pred)
print(f"Recall Rate: {recall:.4f}")
```

False Positive Rate: 0.4336
Cross-Validated Score (Negative MSE): -0.4895
Recall Rate: 0.5393

```
[21]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix,
    ConfusionMatrixDisplay, roc_auc_score, roc_curve
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
```

```

# === STEP 1: Encode categorical variables ===
df['gender_num'] = df['gender'].str.lower().map({'male': 1, 'female': 0})
df['gender_num_o'] = df['gender_o'].str.lower().map({'male': 1, 'female': 0})
df['race_num'] = df['race'].astype('category').cat.codes
df['race_o_num'] = df['race_o'].astype('category').cat.codes

all_fields = pd.concat([df["field"], df["field_o"]], ignore_index=True).
    ↪astype('category')
consistent_categories = all_fields.cat.categories
df["field"] = pd.Categorical(df["field"], categories=consistent_categories)
df["field_o"] = pd.Categorical(df["field_o"], categories=consistent_categories)
df["field_num"] = df["field"].cat.codes
df["field_num_o"] = df["field_o"].cat.codes

# === STEP 2: Define predictors and outcome ===
features = [
    'd_age', 'samerace', 'interests_correlate', 'same_field',
    'bias_attractive', 'bias_sincere', 'bias_intelligence',
    'bias_funny', 'bias_ambition',
    'gender_num', 'race_num', 'age', 'field_num',
    'race_o_num', 'age_o',
    'attractive_partner', 'sincere_partner', 'funny_partner',
    'intelligence_partner', 'ambition_partner', 'gender_num_o',
    'field_num_o', 'attractive', 'sincere', 'intelligence', 'funny',
    'ambition'
]
X = df[features]
y = pd.to_numeric(df['prediction_like'], errors='coerce')

# === STEP 3: Drop missing values ===
model_df = pd.concat([X, y], axis=1).dropna()
X_clean = model_df[features]
y_clean = model_df['prediction_like']

# === STEP 4: Train/test split ===
X_train, X_test, y_train, y_test = train_test_split(X_clean, y_clean,
    ↪test_size=0.2, random_state=42)

# === STEP 5: Logistic Regression Pipeline with Class Balancing ===
logit_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('logreg', LogisticRegression(penalty=None, solver='lbfgs', max_iter=1000,
    ↪class_weight='balanced'))
])
logit_pipeline.fit(X_train, y_train)

```

```

# === STEP 6: Predictions and Probabilities ===
y_pred = logit_pipeline.predict(X_test)
y_prob = logit_pipeline.predict_proba(X_test)[:, 1]

# === STEP 7: Classification Report ===
print("\n=== Logistic Regression Classification Report (Balanced) ===")
print(classification_report(y_test, y_pred))

# === STEP 8: Confusion Matrix ===
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap='Blues')
plt.title("Confusion Matrix - Logistic Regression (Balanced)")
plt.show()

# === STEP 9: ROC Curve & AUC ===
fpr, tpr, _ = roc_curve(y_test, y_prob)
auc = roc_auc_score(y_test, y_prob)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {auc:.2f})', linewidth=2)
plt.plot([0, 1], [0, 1], 'k--', label='Random Guess')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Logistic Regression (Balanced)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

print(f"\nAUC Score: {auc:.4f}")


# === STEP 10: Coefficients Summary ===
logreg_model = logit_pipeline.named_steps['logreg']
coefs = logreg_model.coef_[0]
coef_df = pd.DataFrame({'Feature': features, 'Coefficient': coefs})
print("\n=== Logistic Regression Coefficients (Balanced) ===")
print(coef_df.sort_values(by='Coefficient', key=abs, ascending=False))

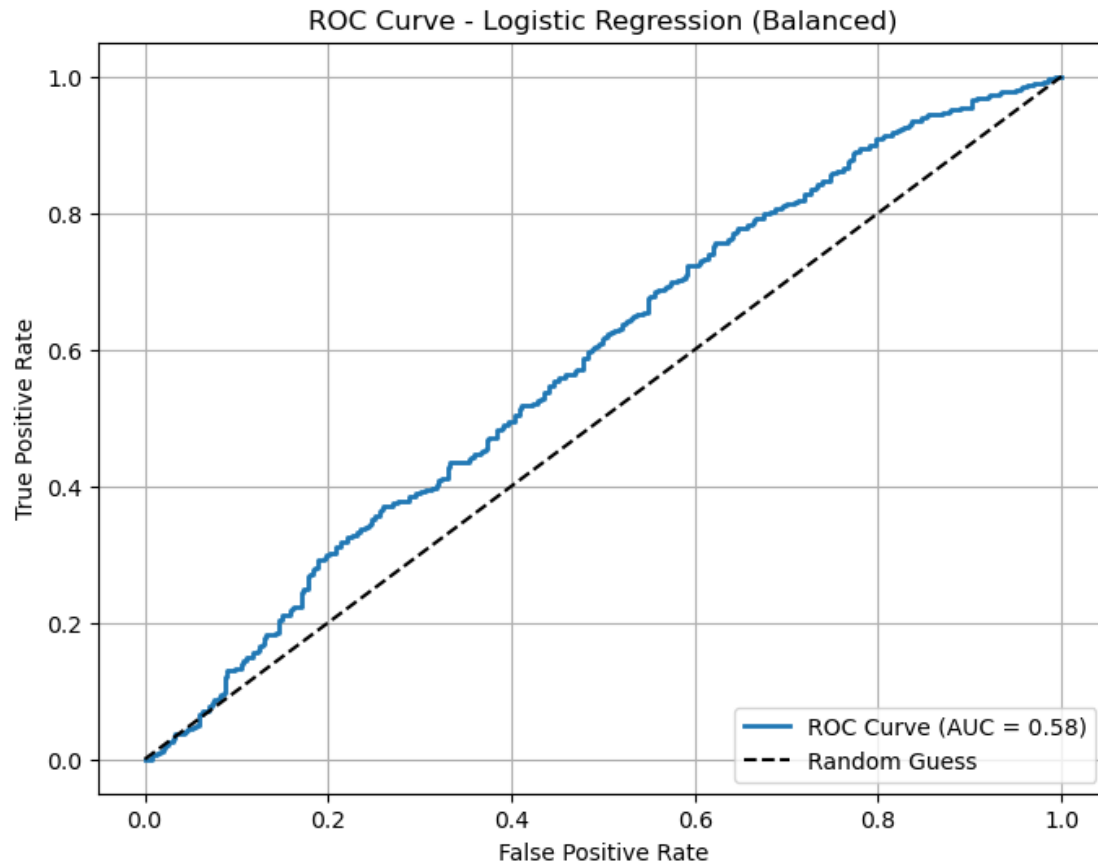
```

=== Logistic Regression Classification Report (Balanced) ===

	precision	recall	f1-score	support
0	0.39	0.56	0.46	399
1	0.70	0.54	0.61	751
accuracy			0.55	1150
macro avg	0.54	0.55	0.53	1150

weighted avg 0.59 0.55 0.56 1150





AUC Score: 0.5816

=== Logistic Regression Coefficients (Balanced) ===

	Feature	Coefficient
6	bias_intelligence	-0.256378
5	bias_sincere	0.198173
23	sincere	-0.172667
24	intelligence	0.166952
7	bias_funny	-0.165158
25	funny	0.149857
8	bias_ambition	0.117470
26	ambition	-0.111364
4	bias_attractive	-0.075939
2	interests_correlate	0.065024
18	intelligence_partner	-0.054669
12	field_num	-0.053585
16	sincere_partner	-0.049315
10	race_num	-0.037101
14	age_o	-0.033461

15	attractive_partner	-0.032829
9	gender_num	-0.029148
20	gender_num_o	0.029148
0	d_age	-0.029006
1	samerace	-0.028844
21	field_num_o	-0.026981
19	ambition_partner	-0.024045
13	race_o_num	-0.013002
3	same_field	-0.010689
22	attractive	-0.004807
11	age	0.001831
17	funny_partner	0.000234

```
[22]: from sklearn.metrics import recall_score
from sklearn.model_selection import cross_val_score

# === False Positive Rate ===
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
false_positive_rate = fp / (fp + tn)
print(f"\nFalse Positive Rate: {false_positive_rate:.4f}")

# === Cross-Validated Score (Negative MSE) ===
neg_mse_scores = cross_val_score(
    logit_pipeline, X_clean, y_clean,
    cv=5, scoring='neg_mean_squared_error'
)
mean_neg_mse = neg_mse_scores.mean()
print(f"Cross-Validated Score (Negative MSE): {mean_neg_mse:.4f}")

# === Recall Rate ===
recall = recall_score(y_test, y_pred)
print(f"Recall Rate: {recall:.4f}")
```

False Positive Rate: 0.4361
Cross-Validated Score (Negative MSE): -0.4897
Recall Rate: 0.5353

```
[23]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
```

```

# Encode categorical variables
df['gender_num'] = df['gender'].str.lower().map({'male': 1, 'female': 0})
df['gender_num_o'] = df['gender_o'].str.lower().map({'male': 1, 'female': 0})
df['race_num'] = df['race'].astype('category').cat.codes
all_fields = pd.concat([df["field"], df["field_o"]], ignore_index=True).
    ↪astype('category')
consistent_categories = all_fields.cat.categories

# Apply the same categories to both columns
df["field"] = pd.Categorical(df["field"], categories=consistent_categories)
df["field_o"] = pd.Categorical(df["field_o"], categories=consistent_categories)

# Now generate codes
df["field_num"] = df["field"].cat.codes
df["field_num_o"] = df["field_o"].cat.codes

# Define predictors and outcome
features = [
    'd_age', 'samerace', 'interests_correlate', 'same_field',
    'bias_attractive', 'bias_sincere', 'bias_intelligence',
    'bias_funny', 'bias_ambition',
    'gender_num', 'race_num', 'age', 'field_num',
    'race_o_num', 'age_o',
    'attractive_partner', 'sincere_partner', 'funny_partner',
    'intelligence_partner', 'ambition_partner', 'gender_num_o',
    'field_num_o', 'attractive', 'sincere', 'intelligence', 'funny',
    'ambition'
]
X = df[features]
y = pd.to_numeric(df['prediction_like'], errors='coerce')

# Drop rows with any missing data
model_df = pd.concat([X, y], axis=1).dropna()
X_clean = model_df[features]
y_clean = model_df['prediction_like']

# Train/test split
X_train, X_test, y_train, y_test = train_test_split(X_clean, y_clean,
    ↪test_size=0.2, random_state=42)

# Loop through different values of C
C_values = [0.001, 0.01, 0.1, 1, 10, 100]
train_acc = []
test_acc = []
coef_norm = []

for C in C_values:

```

```

pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('logreg', LogisticRegression(C=C, penalty='l2', solver='liblinear'))
])

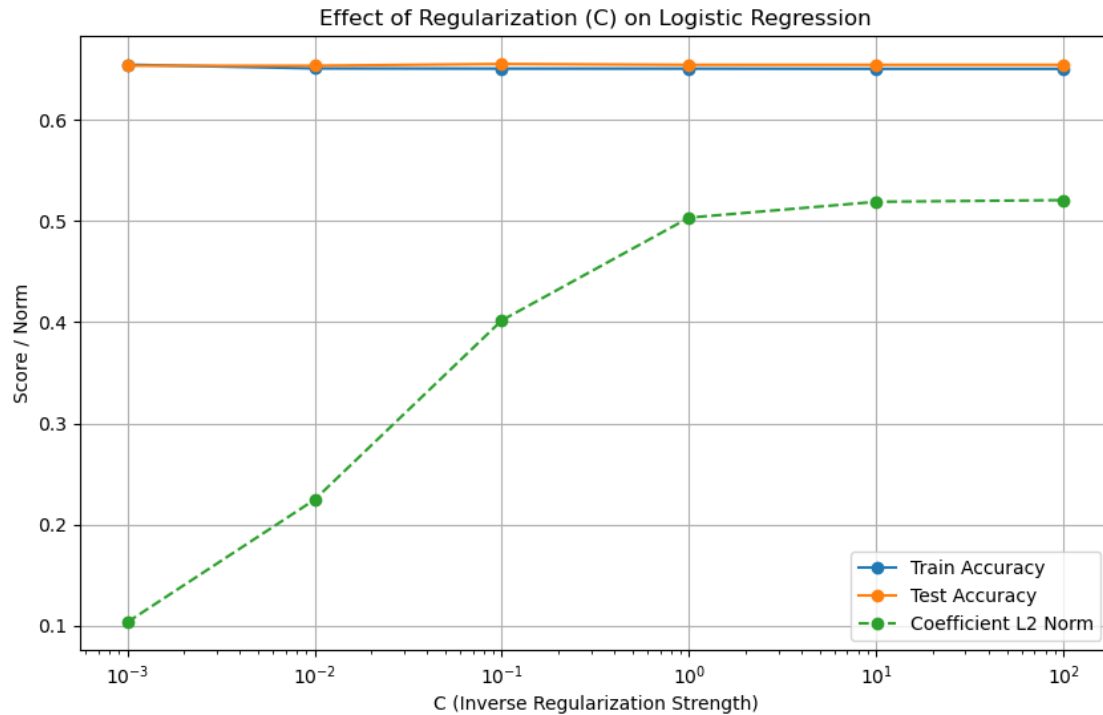
pipeline.fit(X_train, y_train)
y_train_pred = pipeline.predict(X_train)
y_test_pred = pipeline.predict(X_test)

train_acc.append(accuracy_score(y_train, y_train_pred))
test_acc.append(accuracy_score(y_test, y_test_pred))

coefs = pipeline.named_steps['logreg'].coef_
coef_norm.append(np.linalg.norm(coefs)) # L2 norm of coefficients

# Plot performance vs. C
plt.figure(figsize=(10, 6))
plt.plot(C_values, train_acc, label='Train Accuracy', marker='o')
plt.plot(C_values, test_acc, label='Test Accuracy', marker='o')
plt.plot(C_values, coef_norm, label='Coefficient L2 Norm', marker='o', linestyle='--')
plt.xscale('log')
plt.xlabel('C (Inverse Regularization Strength)')
plt.ylabel('Score / Norm')
plt.title('Effect of Regularization (C) on Logistic Regression')
plt.legend()
plt.grid(True)
plt.show()

```



```
[24]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import classification_report, confusion_matrix,
    ConfusionMatrixDisplay, roc_auc_score, roc_curve
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

# === STEP 1: Encode categorical variables ===
df['gender_num'] = df['gender'].str.lower().map({'male': 1, 'female': 0})
df['gender_num_o'] = df['gender_o'].str.lower().map({'male': 1, 'female': 0})
df['race_num'] = df['race'].astype('category').cat.codes
df['race_o_num'] = df['race_o'].astype('category').cat.codes

all_fields = pd.concat([df["field"], df["field_o"]], ignore_index=True).
    astype('category')
consistent_categories = all_fields.cat.categories
df["field"] = pd.Categorical(df["field"], categories=consistent_categories)
df["field_o"] = pd.Categorical(df["field_o"], categories=consistent_categories)
df["field_num"] = df["field"].cat.codes
df["field_num_o"] = df["field_o"].cat.codes
```

```

# === STEP 2: Define predictors ===
features = [
    'd_age', 'samerace', 'interests_correlate', 'same_field',
    'bias_attractive', 'bias_sincere', 'bias_intelligence',
    'bias_funny', 'bias_ambition',
    'gender_num', 'race_num', 'age', 'field_num',
    'race_o_num', 'age_o',
    'attractive_partner', 'sincere_partner', 'funny_partner',
    'intelligence_partner', 'ambition_partner', 'gender_num_o',
    'field_num_o', 'attractive', 'sincere', 'intelligence', 'funny',
    'ambition'
]
X = df[features]
y = pd.to_numeric(df['prediction_like'], errors='coerce')

# === STEP 3: Drop missing values ===
model_df = pd.concat([X, y], axis=1).dropna()
X_clean = model_df[features]
y_clean = model_df['prediction_like']

# === STEP 4: Train-test split ===
X_train, X_test, y_train, y_test = train_test_split(X_clean, y_clean,
    ↪test_size=0.2, random_state=42)

# === STEP 5: Pipeline with class balancing ===
rf_pipe = Pipeline([
    ('scaler', StandardScaler()),
    ('rf', RandomForestClassifier(random_state=42, class_weight='balanced'))
])

# === STEP 6: Grid Search Tuning ===
param_grid = {
    'rf__n_estimators': [100, 200],
    'rf__max_depth': [None, 5, 10],
    'rf__min_samples_split': [2, 5],
    'rf__min_samples_leaf': [1, 2]
}
grid = GridSearchCV(rf_pipe, param_grid, cv=5, scoring='accuracy')
grid.fit(X_train, y_train)

# === STEP 7: Evaluation ===
y_pred = grid.predict(X_test)
y_prob = grid.predict_proba(X_test)[:, 1]

print(f"Best Parameters: {grid.best_params_}")
print("\n=== Classification Report (Balanced Random Forest) ===")
print(classification_report(y_test, y_pred))

```

```

# === STEP 8: Confusion Matrix ===
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap='Blues')
plt.title("Confusion Matrix - Random Forest (Balanced)")
plt.show()

# === STEP 9: ROC Curve and AUC ===
fpr, tpr, _ = roc_curve(y_test, y_prob)
auc_score = roc_auc_score(y_test, y_prob)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {auc_score:.2f})', linewidth=2)
plt.plot([0, 1], [0, 1], 'k--', label='Random Guess')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Random Forest (Balanced)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

print(f"\nAUC Score: {auc_score:.4f}")

# === STEP 10: Feature Importance ===
best_rf = grid.best_estimator_.named_steps['rf']
importances = best_rf.feature_importances_
importance_df = pd.DataFrame({'Feature': features, 'Importance': importances})
importance_df = importance_df.sort_values(by='Importance', ascending=False)

# Print top features
print("\n=== Feature Importances (Balanced RF) ===")
print(importance_df)

# Plot
plt.figure(figsize=(10, 6))
plt.barh(importance_df['Feature'], importance_df['Importance'])
plt.xlabel("Importance Score")
plt.title("Random Forest Feature Importances (Balanced)")
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()

```

```

Best Parameters: {'rf__max_depth': None, 'rf__min_samples_leaf': 1,
'rf__min_samples_split': 2, 'rf__n_estimators': 200}


```


```

=== Classification Report (Balanced Random Forest) ===

```

	precision	recall	f1-score	support
0	0.58	0.24	0.34	399
1	0.69	0.91	0.79	751
accuracy			0.68	1150
macro avg	0.64	0.57	0.56	1150
weighted avg	0.65	0.68	0.63	1150



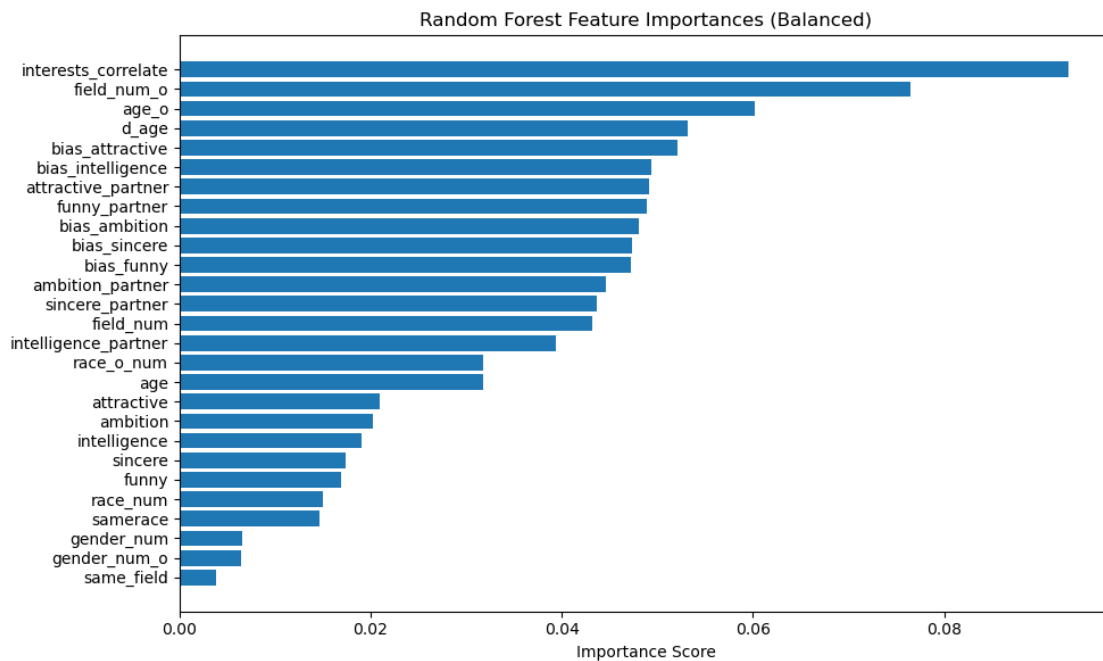


AUC Score: 0.6431

=== Feature Importances (Balanced RF) ===

	Feature	Importance
2	interests_correlate	0.093009
21	field_num_o	0.076490
14	age_o	0.060217
0	d_age	0.053130
4	bias_attractive	0.052048
6	bias_intelligence	0.049330
15	attractive_partner	0.049122
17	funny_partner	0.048944
8	bias_ambition	0.048091
5	bias_sincere	0.047338
7	bias_funny	0.047194
19	ambition_partner	0.044592
16	sincere_partner	0.043707
12	field_num	0.043232
18	intelligence_partner	0.039408

13	race_o_num	0.031732
11	age	0.031721
22	attractive	0.020900
26	ambition	0.020169
24	intelligence	0.018985
23	sincere	0.017391
25	funny	0.016894
10	race_num	0.015024
1	samerace	0.014583
9	gender_num	0.006520
20	gender_num_o	0.006472
3	same_field	0.003758

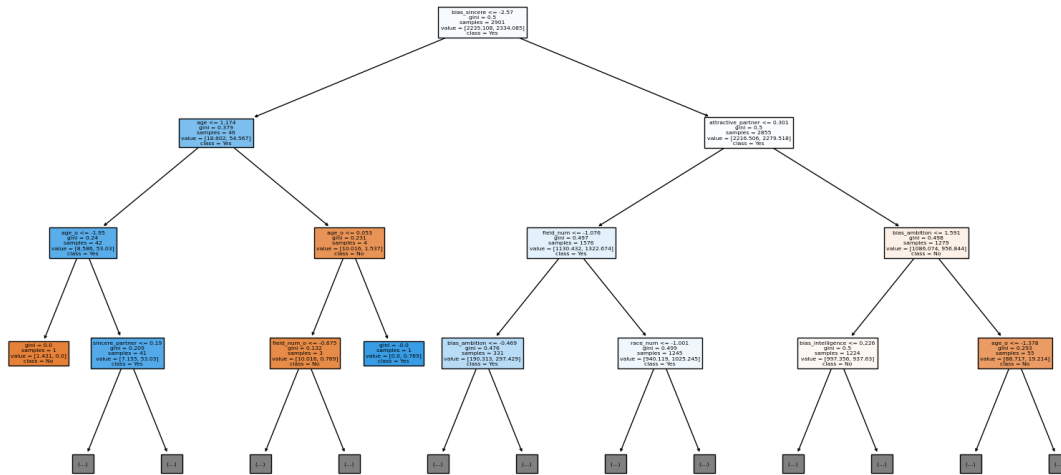


```
[25]: from sklearn.tree import plot_tree, export_text
import matplotlib.pyplot as plt

# === STEP 11: Visualize One Decision Tree ===
# Extract one tree from the best Random Forest model
one_tree = best_rf.estimators_[0] # You can change index to view other trees

# Plot the decision tree
plt.figure(figsize=(20, 10))
plot_tree(one_tree, feature_names=features, class_names=["No", "Yes"],
          filled=True, max_depth=3)
plt.title("Single Decision Tree from Random Forest (Truncated at depth=3)")
plt.show()
```

Single Decision Tree from Random Forest (Truncated at depth=3)



```
[26]: from sklearn.metrics import recall_score
from sklearn.model_selection import cross_val_score

# === False Positive Rate ===
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
false_positive_rate = fp / (fp + tn)
print(f"\nFalse Positive Rate: {false_positive_rate:.4f}")

# === Cross-Validated Score (Negative MSE) ===
neg_mse_scores = cross_val_score(
    grid.best_estimator_, X_clean, y_clean,
    cv=5, scoring='neg_mean_squared_error'
)
mean_neg_mse = neg_mse_scores.mean()
print(f"Cross-Validated Score (Negative MSE): {mean_neg_mse:.4f}")

# === Recall Rate ===
recall = recall_score(y_test, y_pred)
print(f"Recall Rate: {recall:.4f}")
```

False Positive Rate: 0.7644

Cross-Validated Score (Negative MSE): -0.3749

Recall Rate: 0.9095

```
[27]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```

from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import classification_report, roc_auc_score, roc_curve, \
    ↪confusion_matrix, ConfusionMatrixDisplay
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.utils.class_weight import compute_sample_weight

# === STEP 1: Encode categorical variables ===
df['gender_num'] = df['gender'].str.lower().map({'male': 1, 'female': 0})
df['gender_num_o'] = df['gender_o'].str.lower().map({'male': 1, 'female': 0})
df['race_num'] = df['race'].astype('category').cat.codes
df['race_o_num'] = df['race_o'].astype('category').cat.codes

all_fields = pd.concat([df["field"], df["field_o"]], ignore_index=True).
    ↪astype('category')
consistent_categories = all_fields.cat.categories
df["field"] = pd.Categorical(df["field"], categories=consistent_categories)
df["field_o"] = pd.Categorical(df["field_o"], categories=consistent_categories)
df["field_num"] = df["field"].cat.codes
df["field_num_o"] = df["field_o"].cat.codes

# === STEP 2: Define predictors and target ===
features = [
    'd_age', 'samerace', 'interests_correlate', 'same_field',
    'bias_attractive', 'bias_sincere', 'bias_intelligence',
    'bias_funny', 'bias_ambition',
    'gender_num', 'race_num', 'age', 'field_num',
    'race_o_num', 'age_o',
    'attractive_partner', 'sincere_partner', 'funny_partner',
    'intelligence_partner', 'ambition_partner', 'gender_num_o',
    'field_num_o', 'attractive', 'sincere', 'intelligence', 'funny',
    'ambition'
]
X = df[features]
y = pd.to_numeric(df['prediction_like'], errors='coerce')

# === STEP 3: Drop missing values ===
model_df = pd.concat([X, y], axis=1).dropna()
X_clean = model_df[features]
y_clean = model_df['prediction_like']

# === STEP 4: Train/test split ===
X_train, X_test, y_train, y_test = train_test_split(X_clean, y_clean, \
    ↪test_size=0.2, random_state=42)

# === STEP 5: Sample weights for class balancing ===

```

```

sample_weights = compute_sample_weight(class_weight='balanced', y=y_train)

# === STEP 6: Boosting pipeline ===
boost_pipe = Pipeline([
    ('scaler', StandardScaler()),
    ('boost', GradientBoostingClassifier(random_state=42))
])

# === STEP 7: Hyperparameter tuning ===
param_grid = {
    'boost__n_estimators': [50, 100, 150],
    'boost__learning_rate': [0.01, 0.05, 0.1],
    'boost__max_depth': [2, 3, 5]
}

grid = GridSearchCV(boost_pipe, param_grid, cv=5, scoring='accuracy')
grid.fit(X_train, y_train, boost__sample_weight=sample_weights)

# === STEP 8: Predict and evaluate ===
y_pred = grid.predict(X_test)
y_prob = grid.predict_proba(X_test)[:, 1]

print(f"Best parameters: {grid.best_params_}")
print("\n=== Gradient Boosting (Balanced) Classification Report ===")
print(classification_report(y_test, y_pred))

# === Confusion matrix ===
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap='Blues')
plt.title("Confusion Matrix - Gradient Boosting (Balanced)")
plt.show()

# === ROC Curve and AUC ===
fpr, tpr, _ = roc_curve(y_test, y_prob)
auc = roc_auc_score(y_test, y_prob)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {auc:.2f})', linewidth=2)
plt.plot([0, 1], [0, 1], 'k--', label='Random Guess')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Gradient Boosting (Balanced)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()


```

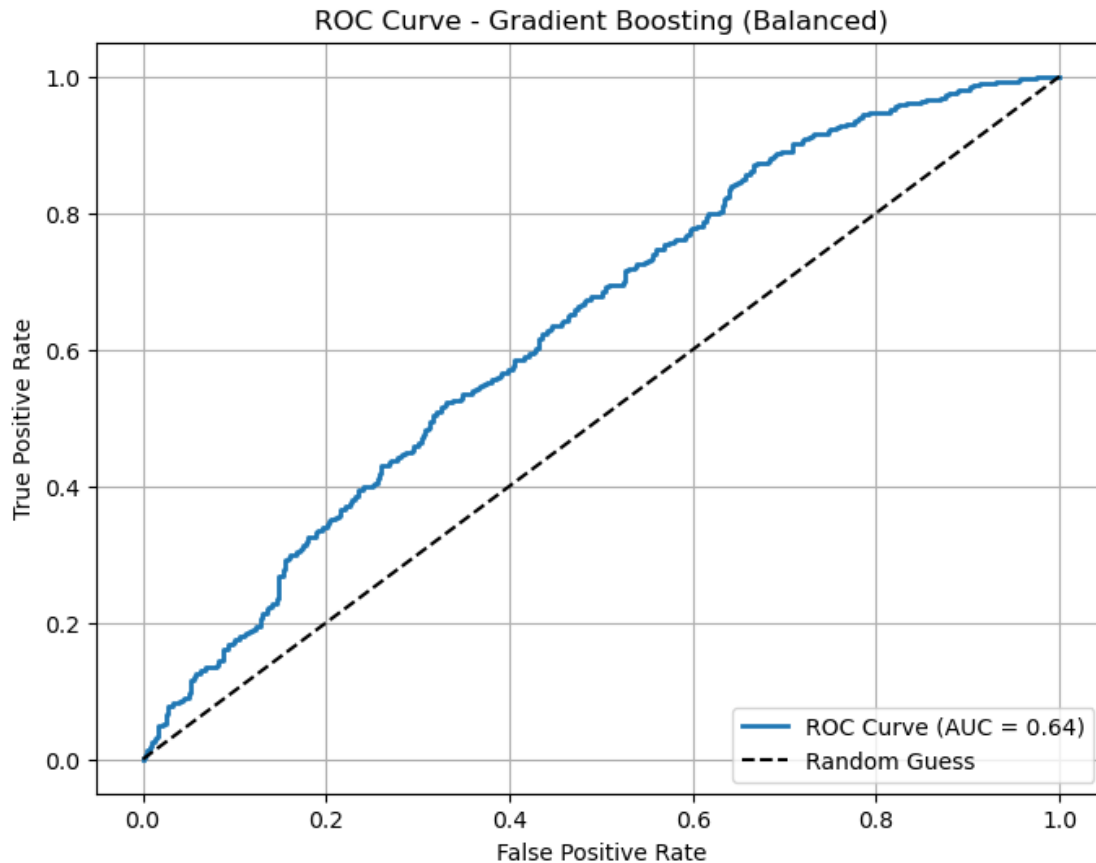
```
print(f"\nAUC Score: {auc:.4f}")
```

Best parameters: {'boost__learning_rate': 0.1, 'boost__max_depth': 5, 'boost__n_estimators': 150}

=== Gradient Boosting (Balanced) Classification Report ===

	precision	recall	f1-score	support
0	0.46	0.47	0.47	399
1	0.72	0.71	0.71	751
accuracy			0.63	1150
macro avg	0.59	0.59	0.59	1150
weighted avg	0.63	0.63	0.63	1150





AUC Score: 0.6364

```
[28]: from sklearn.metrics import recall_score
from sklearn.model_selection import cross_val_score

# === False Positive Rate ===
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
false_positive_rate = fp / (fp + tn)
print(f"\nFalse Positive Rate: {false_positive_rate:.4f}")

# === Cross-Validated Score (Negative MSE) ===
neg_mse_scores = cross_val_score(
    grid.best_estimator_, X_clean, y_clean,
    cv=5, scoring='neg_mean_squared_error'
)
mean_neg_mse = neg_mse_scores.mean()
print(f"Cross-Validated Score (Negative MSE): {mean_neg_mse:.4f}")

# === Recall Rate ===
```

```
recall = recall_score(y_test, y_pred)
print(f"Recall Rate: {recall:.4f}")
```

False Positive Rate: 0.5263

Cross-Validated Score (Negative MSE): -0.4104

Recall Rate: 0.7057

[]: